Approach

The solution follows a simple automation pipeline:

1. Read a meeting transcript from a text file.
2. Use an LLM through openrouter to get output, action items, including title, assignee, due date, and context.
3. Check the output actions with a 2nd AI model to improve JSON file.
4. Automatically create jira tasks for each action item using the jira REST API.
5. Post a summary of all created tickets to slack, including ticket links, assignees, and due dates.

The goal was to minimize manual effort after meetings by turning unstructured conversations into trackable tasks with a single command.

Assumptions

Meeting transcripts are available as plain text. Action items are discussed explicitly enough for an LLM to identify ownership and intent. jira project permissions and Slack bot permissions are already configured. Users prefer task creation automation over manual review before ticket creation. The extracted action items are accurate enough for direct ticket generation, with the second LLM pass acting as a validation step.

Additional Features With More Time

Human approval: Allow users to review, edit, approve, or reject extracted tasks before Jira creation.

Meeting summaries: Generate concise meeting summaries, decisions made, blockers, and risks alongside action items.

Smart assignee mapping: Map participant names to Jira users automatically and assign tickets directly.

Due dates: Infer due dates from phrases such as "by Friday", "next Tuesday", or "within two weeks".

Detect duplicates: Prevent duplicate Jira tickets by checking against existing open issues.

Support many apps: Extend integrations like, Microsoft Teams, Google Meet, Zoom, Notion, Linear, Asana

Future Vision

The current implementation focuses on action-item extraction and execution. With additional development, it could evolve into a complete meeting intelligence platform that automatically captures decisions, tracks ownership, monitors follow-through, and integrates directly into existing team workflows.